# Final Project: Description Using Models

UCB MIDS 203 Fall 2026 Section 6 Final Project Report

Vlad Lee, Sophia Liang, Mark Randall

## 1. Introduction

TODO:
* This is copied directly from the proposal - if we hit our 4 page cap then we may need to come back to these first two paragraphs and condense them

Our research question is to describe at a country level how the Life Expectancy of residents is related to economic development, measured by GDP per capita. Our audience is individuals who are not socio-economic specialists, but would like to expand their worldview of how a wide range of society metrics can influence the typical individual's income, and understand the broader patterns in global development and living standards. The larger context of the question lies in discussions of economic development, global inequality, and public health, where researchers and policymakers aim to understand why countries experience very different health outcomes and living standards. This topic has been widely studied, not least of which by Gapminder itself, which aims to help individuals rid themselves of common global misconceptions. Predicting life expectancy is also studied by Samuel Preston in 1975, who proposed a Preston curve, which demonstrates a concave relationship between life expectancy with GDP per capita. Additionally, Roffia et al. conducted a longitudinal study to determine life expectancy at birth based on certain socioeconomic and demographic factors. Nevertheless, our research brings together multiple dimensions of national development in one accessible country-level framework, providing a more holistic picture of how multiple kinds of determinants interact and influence life expectancy.

## 2. Datasource

TODO:
* should we indicate that each variable was a separate file? And if so, should we write the R code to read the 10 files and combine them into a single data frame?
* note that there are two countries with a missing lex value - I just deleted them as outliers but we should mention this in the data wrangling - Lichtenstein and Holy See

We have sourced our data from www.gapminder.org, a Swedish organization that curates reliable data on global trends, compiled from sources such as the World Bank. The data includes annual country level observations for 193 countries spanning multiple decades. For this project we construct a cross sectional dataset using data from 2022. The data is presented as one row per country (195 observations) and we will use country metrics as our features that are normalized (e.g. per person) to remove country size effects.

### 2.1 Columns

Each country is represented by a 3-letter country code and country long name. The outcome variable for our analysis is life expectancy (lex), which captures overall population health. Our primary explanatory variable is GDP per capita (gdp_pcap), which measures economic development. Though the online dataset includes hundreds of indicators, we have chosen a small subset to represent broad categories of health, inequality, population demographics (age), economy and infrastructure:
* average_daily_income
* broadband_subscribers_per_100_p
* cell_phones_per_100_people
* children_per_woman_total_fertil
* co2_pcap_cons
* gini
* median_age_years
* population_density_per_square_k

## 3. Primary Analysis

TODO:
* add model assumptions
* Polish charts (labels, number formats, titles etc). Add chart numbers?
* review and expand/add a description on each chart
* do we need to describe why we usew heteroskedastic variance in the t-test? The scatter plot of the data shows a relatively consistent variance.
* Review and elaborate on t-test reuslts and effect size interpretation

Our primary anlysis aims to describe the relationship between life expectancy at birth and GDP per capita across our country data in 2022. GDP per capita is a currency measurement and currency data (e.g. salary, earnings) is typically well suited to a logarithmic transformation. We plot the observed life expectancy data against the observed GDP per capita data:

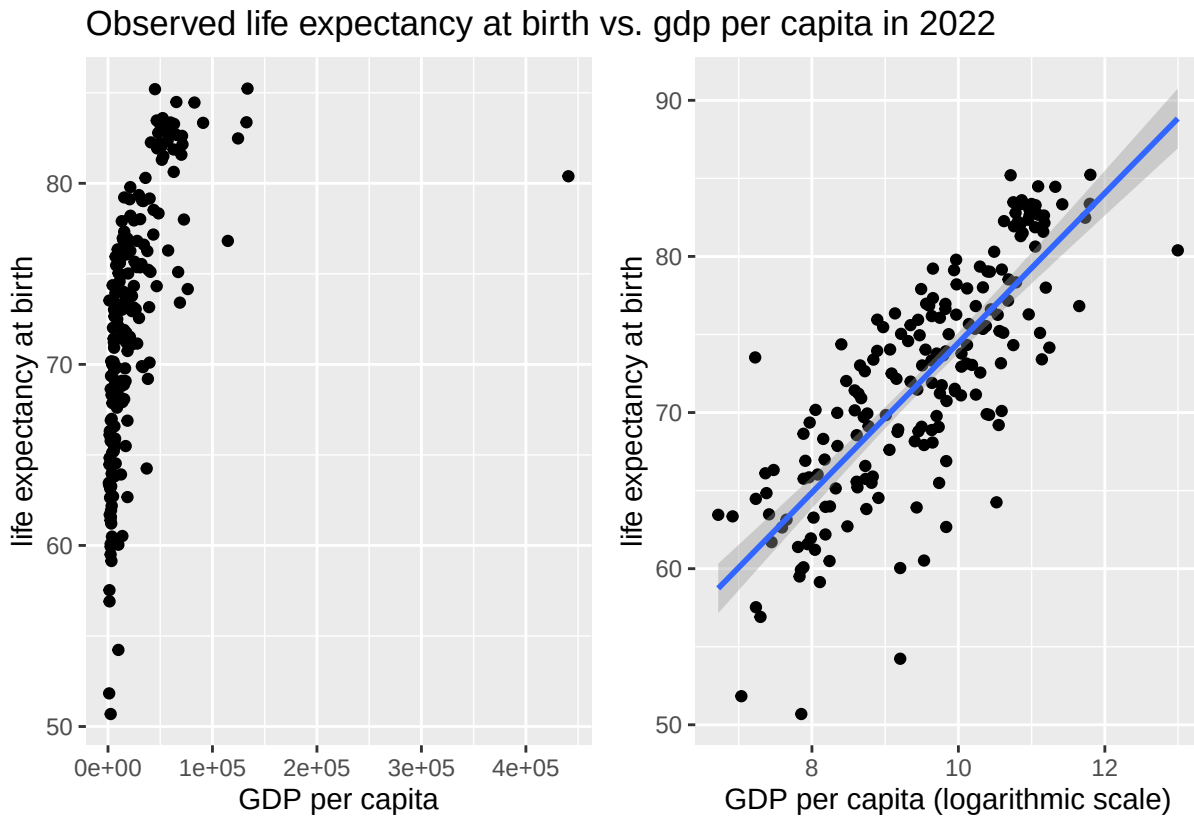


The correlation coefficient of life expectancy against the log of GPD per capita is 0.796, supporting the clear linear relationship indicated in the scatter plot.

We construct our primary regression model as:

$$\text{lex} = \alpha + \beta \cdot \log(\text{GDP per capita})$$

Following the application of OLS, our model is determined as:

$$\text{lex} = 26.501 + 4.796 \cdot \log(\text{GDP per capita})$$

The observed vs. fitted plot suggests that log-GDP-pcap model captures the general positive correlation with life expectancy, since fitted values track the observed values fairly well. However, the residuals vs. fitted values plot shows a slight systematic curvature rather than a random scatter around zero, specifically at the higher fitted values. This demonstrates that GDP per capital does not capture the relationship completely and additional predictors are necessary to improve the model fit.

We perform a t-test of the simple model coefficients using robust variance. In this model and all subsequent specifications, we report robust standard errors to account for potential heteroskedasticity:

```
##
## t test of coefficients:
```

```
##
##                Estimate Std. Error t value  Pr(>|t|)
## (Intercept)   26.50065    2.70927  9.7815 < 2.2e-16 ***
## log(gdp_pcap)  4.79633    0.27998 17.1311 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The p-value is below 0.001 and much smaller than our target level of significance at 0.05. As a result, we can reject our NULL hypothesis and conclude that there is overwhelming evidence that there is a postive relationship between GDP per capita and life expectancy.

A 10 percent increase in GDP per capita is associated with an approximately 0.48 year increase in life expectancy, though this reflects a cross country relationship rather than a causal effect within a given country.

## 4. Expanded Analysis

TODO:
* polish charts - labels, review log transform for income and population density
* check that we have enough text explaining each chart
* not all of the variables have data for all 193 countries - we should mention this? switch warning back on for the r chunk to see which ones.
* co2 per capita and children per woman seem to have a non-linear shape - maybe a diffferent transformation here?
* average_daily_income completely outshines gdp_pcap - although I assume there must be some form of colinearity here? Should we try to prove this and then exclude average_daily_income?
* Consider: convert one of our numerical variables into a categorical variable by grouping? e.g. high/meidum/low - to demonstrate a LM analysis with factors
* Consider: create a model with interaction effects?

To better understand the relationship between life expectancy and broader measures of development, we extend our baseline model by incorporating additional variables capturing demographics, inequality, and population characteristics. Before constructing the expanded model, we note that several variables are strongly correlated with GDP per capita. In particular, average daily income is highly correlated with GDP per capita (correlation $\approx 0.96$), indicating that both variables capture similar aspects of economic development. Including both in the model would introduce multicollinearity and lead to unstable coefficient estimates. For this reason, we exclude average daily income and retain GDP per capita as our primary measure of economic development.

### 4.1 Demographic Factors

We first extend the model by adding fertility (children per woman) and median age. Both variables are statistically significant, with higher fertility associated with lower life expectancy and higher median age associated with higher life expectancy. Including these variables improves model fit and reduces the magnitude of the GDP coefficient, suggesting that part of the relationship between income and life expectancy operates through demographic factors.

$$
\begin{aligned}
\text{lex} = 46.327 &+ 2.415 \cdot \log(\text{GDP per capita}) \\
&- 1.274 \cdot \text{children per woman} \\
&+ 0.205 \cdot \text{median age}
\end{aligned}
$$

The ANOVA test shows that adding the Gini index significantly improves model fit (F = 18.85, p $< 0.0001$), indicating that inequality adds meaningful explanatory power for life expectancy.

### 4.2 Inequality

We next incorporate inequality by adding the Gini index, a measure of wealth inequality. We observe a negative and statistically significant coefficient for the Gini index, indicating that higher inequality is associated with lower life expectancy, holding other factors constant. The coefficient on median age becomes statistically insignificant in this specification, while the effects of GDP per capita and fertility remain statistically significant and consistent in direction.

$$
\begin{aligned}
\text{lex} = 56.893 &+ 2.252 \cdot \log(\text{GDP per capita}) \\
&- 1.541 \cdot \text{children per woman} + 0.136 \cdot \text{median age} \\
&- 0.159 \cdot \text{Gini}
\end{aligned}
$$

The ANOVA test comparing models with and without inequality shows that adding the Gini index significantly improves model fit, with an F statistic of 15.15 and a p-value of 0.0001. We reject the null hypothesis and conclude that including inequality explains significantly more variation in life expectancy.

## Appendix

Data obtained from Gapminder (https://www.gapminder.org/data/).

Baseline model (Life Expectancy ~ Log(GDP per Capita)) - Shows a strong positive relationship while accounting for skewness in GDP per capita, improving interpretability.

Demographics model (adding fertility and median age) - Adds key population characteristics and shows that factors are statistically significant predictors of life expectancy.

Final model (adding inequality) - Including the Gini index shows that it is statistically significant, while median age becomes insignificant.

Table 1: Regression Models of Life Expectancy

| | Baseline | Demographics | Final Model |
|---|---|---|---|
| Intercept | 26.501*** | 46.327*** | 56.893*** |
| | (2.709) | (5.259) | (6.068) |
| Log GDP per Capita | 4.796*** | 2.415*** | 2.252*** |
| | (0.280) | (0.532) | (0.548) |
| Children per Woman | | −1.274** | −1.541*** |
| | | (0.480) | (0.449) |
| Median Age | | 0.205** | 0.136+ |
| | | (0.078) | (0.075) |
| Gini Index | | | −0.159*** |
| | | | (0.047) |
| Num.Obs. | 193 | 193 | 193 |
| R2 | 0.633 | 0.694 | 0.717 |
| R2 Adj. | 0.631 | 0.689 | 0.711 |
| Std.Errors | HC1 | HC1 | HC1 |
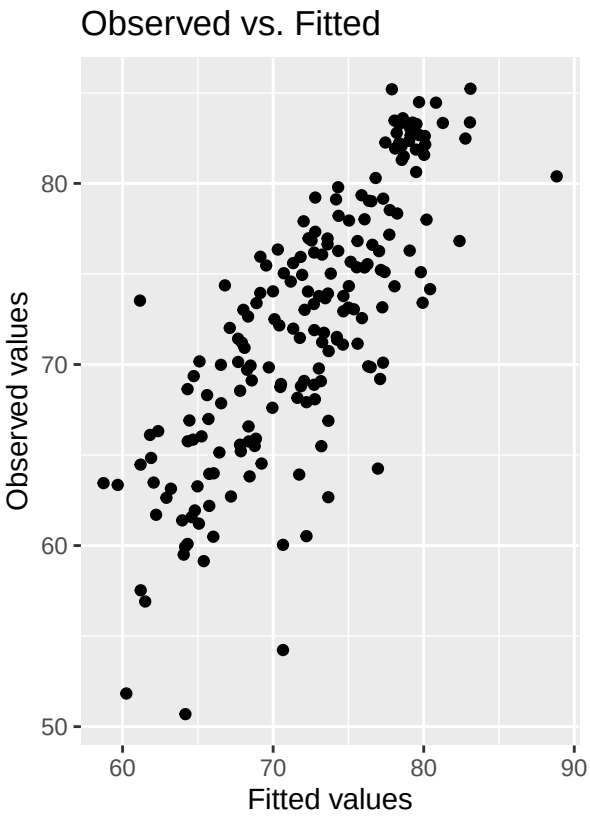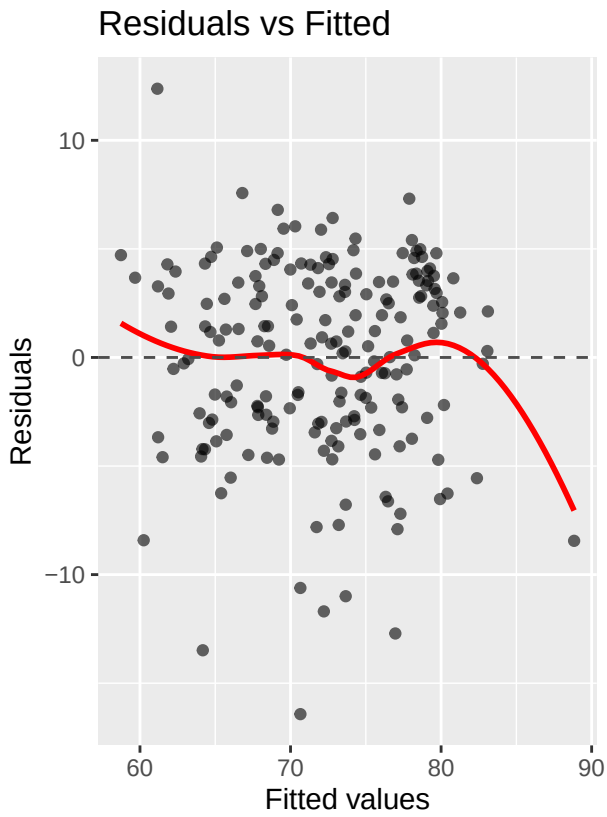
+ p <0.1, * p <0.05, ** p <0.01, *** p <0.001